



Qwen3.8-27B-FP8 inferred attention-core timings

Full Two-tier BF16 Three-tier BF16 Three-tier INT4





TP1 / B1



TP1 / B8



TP4 / B8



Attention core is inferred by subtracting a full-mode dummy-attention run. Different cache-maintenance scaling can bias the residual, especially below 1 ms.